


# Towards Representation Learning with Tractable Probabilistic Models

Antonio Vergari, Nicola Di Mauro and Floriana Esposito

{firstname.lastname@uniba.it}



University of Bari "Aldo Moro", Italy  
Department of Computer Science



LACAM Laboratory  
Machine Learning

## Tractable Probabilistic Models

**Density estimation** is the unsupervised task of learning an estimator for the joint probability distribution  $p(\mathbf{X})$  from a set of i.i.d. samples  $\{\mathbf{x}^i\}_{i=1}^m$  over RVs  $\mathbf{X}$

Given such an estimator, one uses it to answers probabilistic queries about configurations on  $\mathbf{X}$ , i.e. to do **inference**

Different kinds of inference query kinds:


- ⊕  $p(\mathbf{X} = \mathbf{x})$  pointwise evidence (EVI)
- ⊕  $p(\mathbf{E})$ ,  $\mathbf{E} \subset \mathbf{X}$  marginals (MAR)
- ⊕  $p(\mathbf{Q}|\mathbf{E})$ ,  $\mathbf{Q}, \mathbf{E} \subset \mathbf{X}$ ,  $\mathbf{Q} \cap \mathbf{E} = \emptyset$  conditionals (CON)
- ⊕  $\arg \max_{\mathbf{q} \sim \mathbf{Q}} p(\mathbf{q}|\mathbf{E})$  MPE assignments (MPE)
- ⊕  $Z = \sum_{\mathbf{x} \sim \mathbf{X}} \phi(\mathbf{x})$  partition function (Z)
- ⊕ sampling: generate  $\mathbf{x} \sim p(\mathbf{X})$  (SAM)

Good estimators guarantee *exact* and *efficient* inference.

**Tractable Probabilistic Models (TPMs)** are density estimators for which some kind of inference is *tractable*, i.e. polynomial in the number of r.v.s or their domains.


Different TPMs guarantee different inference kinds to be tractable.

**Tree distributions** [6] (\*)




✓ EVI, ✓ SAM, ✓ MAR,  
✓ CON, ✓ MPE, ✓ Z

**d-DNNF** [3] (×)




✓ EVI, ✓ SAM, ✓ MAR,  
✓ CON, ✓ MPE, ✓ Z

**Sum-Product Networks** [7] (×)




✓ EVI, ✓ SAM, ✓ MAR,  
✓ CON, ✓ MPE, ✓ Z

**NADEs** [5] (†)



✓ EVI, ✓ SAM, ✓ MAR,  
✗ CON, ✗ MPE, ✗ Z

**MADEs** [4] (†)



✓ EVI, ✓ SAM, ✓ MAR,  
✗ CON, ✗ MPE, ✗ Z

## Representation Learning with TPMs

**Representation Learning** deals with generating new representations for the initial data  $\{\mathbf{x}^i\}_{i=1}^m$ , e.g. **embeddings**  $\{\mathbf{e}^i | \mathbf{e}^i \in \mathbb{R}^k\}_{i=1}^m$ ,  $k$ -dimensional continuous arrays. Once learned, one can employ them in new tasks such as supervised classification or clustering [2].

Given a TPM  $\theta$  we want to generate an embedding such as:

$$\mathbf{e}^i = f_{p,\theta}(\mathbf{x}^i)$$

for each sample, with  $f$  being a transformation provided by  $\theta$  estimating the probability distribution  $p$ .

Simple idea: exploit the *geometric space* induced by  $p$ , e.g. P-kernels, Fisher vectors... [8], but:

- ⊕ model dependent extraction
- ⊕ closed form analytical derivation needed

We argue that TPMs can be employed as *black box* embedding extractors, on a common ground, by answering generated templated queries.

We define two approaches exploiting a random query generator schema:

I templates are constructed by **random marginal queries**, i.e. generating subsets of r.v.s  $\mathbf{Q}_j \subseteq \mathbf{X}$ ,  $j = 1 \dots d$ :

$$\mathbf{e}_j^i = p_\theta(\mathbf{Q}_j = \mathbf{x}_{\mathbf{Q}_j}^i)$$

II generating a dataset of random *patches* from samples, then training a TPM  $\theta$  on it; embeddings are generated by evaluating  $\theta$  a sliding "window" of size  $d$  along the samples

$$\mathbf{e}_j^i = p_\theta(\mathbf{q}^i), \forall \mathbf{q}^i \sim \mathbf{x}^i, |\mathbf{q}^i| = d$$

Tractable inference is mandatory: for a congruous embedding size  $k$ , one has to perform  $k \cdot m$  queries, e.g. on training split of BMNIST  $k = 10^3$ ,  $m = 5 \cdot 10^4 \rightarrow 5 \cdot 10^7$  evals for approach I

## Experimental Design

Planning an extensive experimentation comprising at least two TPMs from each family class, on a wide range of query types

Empirical evaluation of approach I on five binary image datasets used for classification:

- ⊕ rectangles (REC),  $28 \times 28$  pixels, wide VS tall rectangles
- ⊕ convex (CON),  $28 \times 28$  pixels, convex VS concave shapes
- ⊕ ocr\_letters (OCR),  $16 \times 8$  pixels, ten digits
- ⊕ caltech101 (CAL),  $28 \times 28$  pixels, 101 object shapes
- ⊕ binary MNIST (BMN),  $28 \times 28$  pixels, ten digits

Learning the structure of *differently regularized* TPMs on RVs  $\mathbf{X}$  alone (unsupervised) to compare different *model capacities*:

- ⊕ 3 SPN architectures trained with LearnSPN-b [9] with hyperparameters:  $\rho = 15$  for OCR and  $\rho = 20$  for all the rest,  $m \in \{500, 100, 50\} \rightarrow$  SPN-I, SPN-II, SPN-III), then grid search on  $\alpha \in \{0.1, 0.2, 0.5, 1.0, 2.0\}$
- ⊕ 3 Mixture of trees models (MT) [6] with  $k \in \{3, 15, 30\}$  components ( $\rightarrow$  MT-I, MT-II, MT-III) trained with EM

Extracting embeddings, then training a linear classifier on top of them to predict  $Y$  (supervised):

- ⊕ OVR L2-reg logistic regressor for all representations, grid search for regularization coefficient  $C$  value in  $\{0.0001, 0.001, 0.01, 0.1, 1.0\}$
- ⊕ baseline with initial representations ( $\rightarrow$  LR)

Generating random queries:

- ⊕ up to 1000 randomly generated marginal queries
- ⊕ generating RVs over adjacent pixels in rectangular patches of min sizes of 2, max of 7 pixels for OCR and 10 pixels for the others

## Random query embedding extraction

Approach I demands models for which marginals can be tractable

- ⊕ more flexible: other kinds of inference kinds can be embedded, e.g. more complex query types [1]

Approach II requires only tractable pointwise evidence

- ⊕ largely employable: many more models can answer pointwise queries in a tractable way

**Algorithm 1** randQueryEmbedding( $\mathcal{D}$ ,  $k$ )

```

1: Input: a set of instances  $\mathcal{D} = \{\mathbf{x}^i\}_{i=1}^m$  over r.v.s
    $\mathbf{X} = \{X_1, \dots, X_n\}$ ,  $k$  as the number of features to generate
2: Output: a set of embeddings  $\mathcal{E} = \{\mathbf{e}^i\}_{i=1}^m$ ,  $\mathbf{e}^i \in \mathbb{R}^k$ 
3:  $\theta \leftarrow \text{learnDensityEstimator}(\mathcal{D})$ 
4:  $\mathcal{E} \leftarrow \{\}$ 
5: for  $j = 1, \dots, k$  do
6:    $\mathbf{Q}_j \leftarrow \text{selectRandomRVs}(\mathbf{X})$ 
7:   for  $i = 1, \dots, m$  do
8:      $\mathbf{e}_j^i = p_\theta(\mathbf{x}_{\mathbf{Q}_j}^i)$ 
9:    $\mathcal{E} \leftarrow \mathcal{E} \cup \{\mathbf{e}^i\}$ 
return  $\mathcal{E}$ 

```

**Algorithm 1** randPatchEmbedding( $\mathcal{D}$ ,  $s$ ,  $d$ )

```

1: Input: a set of instances  $\mathcal{D} = \{\mathbf{x}^i\}_{i=1}^m$  over r.v.s  $\mathbf{X} = \{X_1, \dots, X_n\}$ ,  $s$ 
   as the number of patches to extract,  $d$  as the patch length,
2: Output: a set of embeddings  $\mathcal{E} = \{\mathbf{e}^i\}_{i=1}^m$ ,  $\mathbf{e}^i \in \mathbb{R}^k$ 
3:  $\mathcal{R} \leftarrow \{\}$ 
4: for  $i = 1, \dots, s$  do
5:    $\mathbf{x}^{\text{and}} \leftarrow \text{selectRandomSample}(\mathcal{D})$ 
6:    $\mathbf{r}^i \leftarrow \text{extractRandomPatch}(\mathbf{x}^{\text{and}}, d)$ 
7:    $\mathcal{R} \leftarrow \mathcal{R} \cup \{\mathbf{r}^i\}$ 
8:  $\theta \leftarrow \text{learnDensityEstimator}(\mathcal{R})$ 
9:  $\mathcal{E} \leftarrow \{\}$ 
10: for  $i = 1, \dots, m$  do
11:    $j \leftarrow 0$ 
12:   for each patch  $\mathbf{q}^i, |\mathbf{q}^i| = d$  in  $\mathbf{x}^i$  do
13:      $\mathbf{e}_j^i = p_\theta(\mathbf{q}^i)$ 
14:      $j \leftarrow j + 1$ 
15:    $\mathcal{E} \leftarrow \mathcal{E} \cup \{\mathbf{e}^i\}$ 
return  $\mathcal{E}$ 

```


## Results

For all models, less than 300 features to beat the LR baseline  
 $\rightarrow$  the geometric space induced is meaningful/useful


SPN embeddings outperform MT ones all datasets, but CAL

$\rightarrow$  MT scoring best likelihoods on  $\mathbf{X}$  but worst accuracies for  $Y$


More regularized models perform better than specialized ones  
 $\rightarrow$




REC




CON



OCR



CAL



BMN

## References

- [1] Jessa Bekker et al. "Tractable Learning for Complex Probability Queries". In: *Advances in Neural Information Processing Systems 28 (NIPS)*, 2015.
- [2] Yoshua Bengio, Aaron C. Courville, and Pascal Vincent. "Unsupervised Feature Learning and Deep Learning: A Review and New Perspectives". In: *CoRR* abs/1206.5538 (2012).
- [3] Adnan Darwiche. *Modeling and Reasoning with Bayesian Networks*. Cambridge, 2009.
- [4] Mathieu Germain et al. "MADE: Masked Autoencoder for Distribution Estimation". In: *CoRR* abs/1502.03509 (2015).
- [5] Hugo Larochelle and Iain Murray. "The Neural Autoregressive Distribution Estimator". In: *International Conference on Artificial Intelligence and Statistics*, 2011, pp. 29–37.
- [6] Marina Meilă and Michael I. Jordan. "Learning with mixtures of trees". In: *Journal of Machine Learning Research* 1 (2000), pp. 1–48.
- [7] Hoifung Poon and Pedro Domingos. "Sum-Product Networks: a New Deep Architecture". In: *UAI* 2011 (2011).
- [8] John Shawe-Taylor and Nello Cristianini. *Kernel Methods for Pattern Analysis*. New York, NY, USA: Cambridge University Press, 2004. ISBN: 0521813972.
- [9] Antonio Vergari, Nicola Di Mauro, and Floriana Esposito. "Simplifying, Regularizing and Strengthening Sum-Product Network Structure Learning". In: *ECML-PKDD 2015*, 2015.

## Links

paper available at:

<https://arxiv.org/abs/1608.02341>

code available at:

<https://github.com/arranger1044/spyn-repr>

more references:

<https://github.com/arranger1044/awesome-spn>

ECML-PKDD 2016 - 19th-23rd September 2016, Riva del Garda, Italy